Carryover effects in free recall reveal how prior experiences influence memories of new experiences

Jeremy R. Manning[1,*], Kirsten Ziman[1,2], Emily Whitaker[1],
Paxton C. Fitzpatrick[1], Madeline R. Lee[1], Allison M Frantz[1],
Bryan J. Bollinger[1], and Andrew C. Heusser[1,3]

[1]Dartmouth College

[2]Princeton University

[3]Akili Interactive

[*]Corresponding author: jeremy.r.manning@dartmouth.edu

Abstract

We perceive, interpet, and remember ongoing experiences through the lens of our prior experiences. Inferring that we are one type of situation versus another can lead us to interpret the same physical experience differently. In turn, this can affect how we focus our attention, form expectations of what will happen next, remember what is happening now, draw on our prior related experiences, and so on. To study these phenomena, we asked participants to perform simple word list learning tasks. Across different experimental conditions, we held the set of to-be-learned words constant, but we manipulated the orders in which the words were studied. We found that these order manipulations affected not only how the participants recalled the ordered lists, but also how they recalled later randomly ordered lists. Our work shows how structure in our ongoing experiences can exert influence on how we remember unrelated subsequent experiences.

Introduction

Experience is subjective: different people who encounter identical physical experiences can take away very different meanings and memories. One reason is that our subjective experiences in the moment are shaped in part the idiosyncratic prior experiences, memories, goals, thoughts, expectations, and emotions that we bring with us into the present moment. These factors collectively define a *context* for our experiences[?] . situation models: forming expectations, predicting ambiguous future experiences The contexts we encounter help us to construct *situation models*[??] or *schemas*[??] that describe how experiences are likely to unfold based on our prior experiences with similar contextual cues. For example, when we enter a sit-down restaurant, we might expect to be seated at a table, given a menu, and served food. Priming someone to expect a particular situation or context can also influence how they resolve potentail ambiguities in their ongoing experiences, including ambiguous movies and narratives[?] .

Our understanding of how we form situation models and schemas, and how they interact with our subjective experiences and memories, is constrained in part by substantial differences in how we study these processes. Situation models and schemas are most often studied using "naturalistic" stimuli such as narratives and movies[???] . In contrast, our understanding of how we organize our memories has been most widely studied using more traditional paradigms like free recall of random word lists[?] . In free recall, participants study lists of items and are instructed to recall the items in any order they choose. The orders in which words come to mind can provide insights into how participants have organized their memories of the studied words. Because random word lists are unstructured by design, it is not clear if or how non-trivial situation models might apply to these stimuli. Nevertheless, there are *some* commonalities between memory for word lists and memory for real-world experiences.

Like remembering real-world experiences, remembering words on a studied list requires distinguishing the current list from the rest of one's experience. To model this fundamental memory capability, cognitive scientists have posited the existence of a special representation, called *context*, that is associated with each list. According to early theories e.g.[??] context representations are composed of many features which fluctuate from moment to moment, slowly drifting through a multidimensional feature space. During recall, this representation forms part of the retrieval cue, enabling us to distinguish list items from non-list items. Understanding the role of context in memory processes is particularly important in self-cued memory tasks, such as *free recall*, where the retrieval cue is "context" itself.

Over the past half-century, context-based models have enjoyed impressive success at explaining many stereotyped behaviors observed during free recall and other list-learning tasks[??????????]. These phenomena include the well-known recency and primacy effects (superior recall of items from the end and, to a lesser extent, from the beginning of the study list), as well as semantic and temporal clustering effects[?]. The contiguity effect is an example of temporal clustering, which is perhaps the dominant form of organization in free recall. This effect can be seen in the tendency for people to successively recall items that occupied neighboring positions in the study list. For example, if a list contained the sub-sequence "ABSENCE HOLLOW PUPIL" and the participant recalls the word "HOLLOW", it is far more likely that the next response will be either "PUPIL" or "ABSENCE" than some other list item[?]. In addition, there is a strong forward bias in the contiguity effect: subjects make forward transitions (i.e., "HOLLOW" followed by "PUPIL") about twice as often as they make backward transitions, despite an overall tendency to begin recall at the end of the list. There are also striking effects of semantic clustering[?????], whereby the recall of a given item is more likely to be followed by recall of a similar or related item than

a dissimilar or unrelated one. In general, people organize memories for words along a wide variety of stimulus dimensions. As captured by models like the *Context Maintenance and Retrieval Model*[?], the stimulus features associated with each word (e.g. the word's meaning, font size, font color, location on the screen, size of the object the word represents, etc.) are incorporated into the participant's mental context representation[?????]. During a memory test, any of these features may serve as a memory cue, which in turn leads the participant to recall in succession words that share stimulus features.

A key mystery is whether the sorts of situation models and schemas that people use to organize their memories of real-world experiences might map onto the clustering effects that reflect how people organize their memories for word lists. On one hand, situation models and clustering effects both reflect statistical regularities in ongoing experience. Our memory systems exploit these regularities when generating inferences about the unobserved past and yet-to-be-experienced future[?????]. On the other hand, the rich structure of real-world experiences and other naturalistic stimuli that enable people to form deep and meaningful situation models and schemas have no obvious analog in simple word lists. Often lists in free recall studies are explicitly *designed* to be devoid of exploitable temporal structure, for example by sorting the words in a random order[?].

We designed an experimental paradigm to explore how people organize their memories for simple stimuli (word lists) whose temporal properties change across different "situations," analogous to how the content of real-world experiences change across different real-world situations. We asked participants to study and freely recall a series of word lists (Fig. **??**). Across the different conditions in the experiment, we varied the lists' presentation orders in different ways across lists. The studied items (words) were designed to vary along three general dimensions: semantic (word *category*, and physical *size* of the referent), lexicographic (word *length* and *first letter*), and visual (font *color* and

the onscreen *location* of each word). In our main manipulation conditions, we asked participants to study and recall eight lists whose items were sorted by a target feature (e.g., word category). Next, we asked them to study and recall an additional eight lists whose items had the same features, but that were sorted in a random temporal order. We were interested in how these order manipulations affected participants' recall behaviors on early (sorted) lists, as well as how order manipulations on early lists affected recall behaviors on later (unsorted) lists. We used a series of control conditions as a baseline; in these control conditions all of the lists were sorted randomly, but we manipulated the presence or absence of the visual features. Finally, in an *adaptive* experimental condition we used participants' recall behaviors on early lists to manipulate, in real-time, the presentation orders of subsequent lists. In this adaptive condition, we sought to identify potential commonalities within and across participants in how people organized their memories and how those organizational tendancies affect overall performance.

Materials and methods

Participants

We enrolled a total of 491 Dartmouth undergraduate students across 11 experimental conditions. The conditions included two primary controls (feature rich, reduced), two secondary controls (reduced (early), reduced (late)), six order manipulation conditions (category, size, length, first letter, color, and location), and a final adaptive condition. Each of these conditions are described in the *Experimental design* subsection below.

Participants received course credit for enrolling in our study. We asked each participant to fill out a demographic survey that included information about their self-reported age, gender, ethnicity, race, education, vision, reading impairments, medications or recent

injuries, coffee consumption on the day of testing, and level of alertness at the time of testing. All components of the demographics survey were optional. One participant elected not to fill out any part of the demographic survey, and all other participants report some or all of their requested demographic information.

We aimed to run (to completion) at least 60 participants in each of the two primary control conditions and in the adaptive condition. In all other conditions we set a target enrollment of at least 30 participants. Because our data collection efforts were coordinated 12 researchers and multiple testing rooms and computers, it was not feasible for individual experimenters to know how many participants had been run in each experimental condition until the relevant databases were synchronized at the end of each working day. We also over-enrolled participants for each condition to help ensure that we met our minimum enrollment targets even if some participants dropped out of the study prematurely or did not show up for their testing session. This led us to exceed our target enrollments for several conditions.

Participants were assigned to experimental conditions based loosely on their date of participation. (This aspect of our procedure helped us to more easily synchronize the experiment databases across multiple testing computers.) Of the 490 participants who opted to fill out the demographics survey, reported ages ranged from 17 to 31 years (mean: 19.1; standard deviation: 1.356). A total of 318 participants reported their gender as female, 170 as male, and 2 participants declined to report their gender. A total of 442 participants reported their ethnicity as "not Hispanic or Latino," 39 as "Hispanic or Latino," and 9 declined to report their ethnicity. Participants reported their races as White (345 participants), Asian (120 participants), Black or African American (31 participants), American Indian or Alaska Native (11 particiapnts), Native Hawaiian or Other Pacific Islander (4 participants), Mixed race (3 participants), Middle Eastern (1 participant), and Arab (1

participant). A total of 5 participants declined to report their race. We note that several participants reported more than one of racial category. Participants reported their highest degrees achieved as "Some college" (359 participants), "High school graduate" (117 participants), "College graduate" (7 participants), "Some high school" (5 participants), "Doctorate" (1 participant), and "Master's degree" (1 participant). A total of 482 participants reported no reading impairments, and 8 reported mild reading impairments such as mild dyslexia. A total of 489 participants reported having normal color vision and 1 participant reported that they were color blind. A total of 482 participants reported taking no prescription medications and having no recent injuries; 4 participants reported having ADHD, 1 reported having dyslexia, 1 reported having allergies, 1 reported a recently torn ACL/MCL, and 1 reported a concussion from several months prior. The participants reported consuming 0 – 3 cups of coffee prior to the testing session (mean: 0.32 cups; standard deviation: 0.58 cups). Participants reported their current level of alertness, and we converted their responses to numerical scores as follows: "very sluggish" (-2), "a little sluggish" (-1), "neutral" (0), "a little alert" (1), and "very alert" (2). Across all participants, the full range of alertness levels were reported (range: -2 – 2; mean: 0.35; standard deviation: 0.89).

We dropped from our dataset the 1 participant who reported abnormal color vision, as well as 39 participants whose data were corrupted due to technical failures while running the experiment or during the daily database merges. In total, this left usable data from 452 participants, broken down by experimental condition as follows: feature rich (67 participants), reduced (61 participants), reduced (late) (41 participants), reduced (early), (42 participants), category (30 participants), size (30 participants), length (30 participants), first letter (30 participants), color (31 participants), location (30 participants), and adaptive (60 participants). The participant who declined to fill out their demographic survey

participated in the location condition, and we verified verbally that they had normal color vision.

Experimental design

Our experiment is a variant of the classic free recall paradigm that we term *feature-rich free recall*. In feature-rich free recall, participants study 16 lists, each comprised of 16 words that vary along a number of stimulus dimensions (Fig. **??**). The stimulus dimensions include two semantic features related to the *meanings* of the words (semantic category, referrent object size), two lexicographic features related to the *letters* that make up the words (word length in number of letters, identity of the word's first letter), and two visual features that are independent of the words themselves (text color, presentation location). Each list contains four words from each of four different semantic categories and two object sizes; all other stimulus features are randomized. After studying each list, the participant attempts to recall as many words as they can from that list, in any order they choose. Because each individual word is associated with several well-defined (and quantifiable) features, and because each list incorporates a diverse mix of feature values along each dimension, this allows us to evaluate participants' memory fingerprints in rich detail.

Stimuli

Stimuli in our paradigm were 256 English words selected in a previous study[?]. The words all referred to concrete nouns, and were chosen from 15 unique semantic categories: body parts, building-related, cities, clothing, countries, flowers, fruits, insects, instruments, kitchen-related, mammals, (US) states, tools, trees, and vegetables. We also tagged each word according to the approximate size of the object the word referred to. Words were labeled as "small" if the corresponding object was likely able to "fit in a standard shoebox"

Figure 1: Feature-rich free recall. After studying lists comprised of words that vary along several feature dimensions, participants verbally recall words in any order (microphone icon). Each experimental condition manipulates word features and/or presentation orders within and/or across lists. The rows display representative (illustrated) examples of the first lists participants might encounter in each condition. The rectangles during the "Presentation phase" show illustrated screen captures during a series of word presentations. Each word appeared onscreen for 2 seconds, followed by 2 seconds of blank screen. The red microphone icons during the "Recall" phase denote the one minute verbal recall interval. The labels on the right (and corresponding groupings on the left) denote experimental condition labels.

or "large" if the object was larger than a shoebox. Semantic categories varied in how many object sizes they reflected (mean number of different sizes per category: 1.33; standard deviation: 0.49). The numbers of words in each semantic category also varied from 12 – 28 (mean number of words per category: 17.07; standard deviation number of words: 4.65). We also identified lexicographic features for each word, including the words' first letters and lengths (i.e., number of letters). Across all categories, all possible first letters were represented except for 'Q' (average number of unique first letters per category: 11; standard deviation: 2 letters). Word lengths ranged from 3 – 12 letters (average: 6.17 letters; standard deviation: 2.06 letters).

We assigned the categorized words into a total of 16 lists with several constraints. First, we required that each list contained words from exactly 4 unique categories, each with exactly 4 examplars from each category. Second, we required that (across all words on the list) at least one instance of both object sizes were represented. On average, each category was represented in 4.27 lists (standard deviation: 1.16 lists). Aside from these two constraints, we assigned each word to a unique list. After random assignment, each list contained words with an average of 11.13 unique starting letters (standard deviation: 1.15 letters) and an average word length of 6.17 letters (standard deviation: 0.34 letters).

The above assignments of words to lists was performed once across all participants, such that every participant studied the same set of 16 lists. In every condition we randomized the study order of these lists across participants. For participants in some conditions, on some lists, we also randomly varied two additional visual features to each word: the presentation font color, and the word's onscreen location. These attributes were assigned independently for word (and for every participant) at the times the words were displayed onscreen. These visual features were varied for words in all lists and conditions except for the "reduced" condition (all lists), the first eight lists of the "reduced (early)" condition,

and the last eight lists of the "reduced (late)" condition. In these latter cases, words were all presented in black at the center of the experimental computer's display.

To assign a random font color to each word, we selected three integers uniformly and at random between 0 and 255, corresponding to the red (r), green (g), and blue (b) color channels for that word. To assign random presentation locations to each word, we selected two floating point numbers uniformly at random (one for the word's horizontal x coordinate and the other for its vertical y coordinate). The bounds of these coordinates were selected to cover the entire visible area of the display without cutting off any part of the words. The words were shown on 27 in (diagonal) Retina 5K iMac displays (resolution: 5120×2880 pixels).

Most of the experimental manipulations we carried out entailed presenting or sorting the presented words differently on the first eight lists participants studied (which we call *early* lists) versus on the final eight lists they studied (*late* lists). Since every participant studied exactly 16 lists, using this terminology every list was either "early" or "late" depending on its order in the list study sequence.

Real-time speech-to-text processing

Our experimental paradigm incorporates the Google Cloud Speech API speech-to-text engine[?] to automatically transcribe participants' verbal recalls into text. This allows recalls to be transcribed in real time– a distinguishing feature of the experiment; in typical verbal recall experiments the audio data must be parsed manually. In prior work, we used a similar experimental setup (equivalent to the "reduced" condition in the present study) to verify that the automatically transcribed recalls were sufficiently close to human-transcribed recalls to yield reliable data[?] . This real-time speech processing component of the paradigm plays an important role in the "adaptive" condition of the experiment, as

described below.

Random conditions (Fig. ??, top four rows)

We used four "control" conditions to evaluate and explore participants' baseline behaviors. We also used performance on these control conditions to help interpret performance in other "manipulation" conditions. Two control conditions served as "anchorpoints." In the first anchorpoint condition, which we call the *feature rich* condition, we randomly shuffled the presentation order (independently for each participant) of the words on each list. In the second anchorpoint condition, which we call the *reduced* condition, we randomized word presentations as in the feature rich condition. However, rather than assigning each word a random color and location, we instead displayed all of the words in black and at the center of the screen.

In the *reduced (early)* condition, we followed the "reduced" procedure (presenting each word in black at the center of the screen) for early lists, and followed the "feature rich" procedure (presenting each word in a random color and location) for late lists. Finally, in the *reduced (late)* condition, we followed the feature rich procedure for earlylists and the reduced procedure for late lists.

Order manipulation conditions (Fig. ??, middle six rows)

Each of six *order manipulation* conditions used a different feature-based sorting procedure to order words on early lists, where each sorting procedure relied on one relevant feature dimension. All of the irrelevant features varied freely across words on early lists, in that we did not consider irrelevant features in ordering the early lists. However, some features were correlated– for example, some semantic categories of words referred to objects that tended to be a particular size, which means that category and size are not

fully independent. On late lists, the words were always presented in a randomized order (chosen anew for each participant). In all of the order manipulation conditions, we varied words' font colors and onscreen locations, as in the feature rich condition.

Defining feature-based distances. Sorting words according to a given relevant feature requires first defining a distance function for quantifying the dissimilarity between each pair of features. This function varied according to the type of features. Semantic features (category and size) are *categorical*. For these features, we defined a binary distance function: two words were considered to "match" (i.e., have a distance of 0) if their labels are the same (i.e., both from the same semantic category or both of the same size). If two words' labels were different for a given feature, we defined the words to have a distance of 1 for that feature. Lexicographic features (length and first letter) are *discrete*. For these features we defined a discrete distance function. Specifically, we defined the distance between two words as either the absolute difference between their lengths, or the absolute distance between their starting letters in the English alphabet, respectively. For example, two words that started with the same letter would have a "first letter" distance of 0, and words starting with 'J' and 'A' would have a first letter distance of 9. Because words' lengths and letters' positions in the alphabet are always integers, these discrete distances always take on integer values. Finally, the visual features (color and location) are *continuous* and *multivariate*, in that each "feature" takes on multiple (positive) real values. We defined the "color" and "location" distances between two words as the Euclidean distances between their (r, g, b) color or (x, y) location vectors, respectively. Therefore the color and location distance measures always take on positive real values (upper bounded at 441.67 for color, or 27 in for location, reflecting the distances between the corresponding maximally different vectors).

Constructing feature-sorted lists. Given a list of words, a relevant feature, and each word's value(s) for that feature, we developed a stochastic algorithm for (noisily) sorting the words. First, we choose a word uniformly at random from the set of candidates. Next, we compute the distances between the chosen word's feature(s) and the corresponding feature(s) of all yet-to-be-presented words. Third, we convert these distances (between the previously presented word's feature values, a, and the candidate word's feature values, b) to similarity scores:

$$\text{similarity}(a, b) = \exp\{-\tau \cdot \text{distance}(a, b)\}, \tag{1}$$

where $\tau = 1$ in our implementation. We note that increasing the value of τ would amplify the influence of similarity on order, and decreasing the value of τ would diminish the influence of similarity on order. Also note that this approach requires $\tau > 0$. Finally, we computed a set of normalized similarity values by dividing the similarities by their sum:

$$\text{similarity}_{\text{normalized}}(a, b) = \frac{\text{similarity}(a, b)}{\sum_{i=1}^{n} \text{similarity}(a, i)}, \tag{2}$$

where in the demoniator, i takes on each of the n feature values of the to-be-presented words. The resulting set of normalized similarity scores sums to one.

As illustrated in Figure **??**, we use these normalized similarity scores to construct a sequence of "sticks" that we lay end to end in a line. Each of the n sticks corresponds to a single to-be-presented word, and the stick lengths are proportional to the relative similarities between each word's feature value(s) and the feature value(s) of the just-presented word. We choose the next to-be-presented word by moving an indicator along the set of sticks, by a distance chosen uniformly at random on the interval $[0, 1]$. We select the word associated with the stick lying next to the indicator to be presented next. This process continues iteratively (re-computing the similarity scores and stochastically

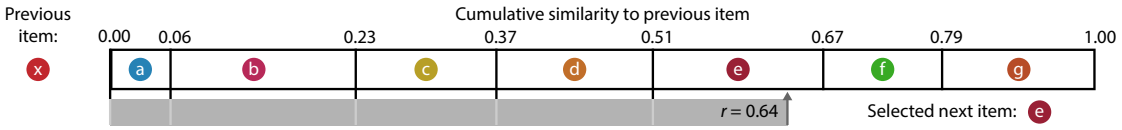

Figure 2: Generating stochastic feature-sorted lists. For a given feature dimension (e.g., color), we compute the similarity (Eqn. **??**) between the feature value(s) of the previous item, x, and all yet-to-be-presented items (a – g). Next, we normalize these similarity scores so that they sum to one. We lay in sequence a set of "sticks," one for each candidate item, whose lengths are equal to these normalized similarity scores. Note that the combined lengths of these sticks is one. To select the next to-be-presented item, we draw a random number, r, from the uniform distibution bounded between 0 and 1 (inclusive). The identity of the next item is given by the stick adjacent to an indicator that moves distance r (starting from 0) along the sequence of sticks. In this case, the next to-be-presented item is e. Note that each item's chances of selection is proportional to its similarity to the previous item, along the given feature dimension.

choosing the next to-be-presented word using the just-presented word) until all of the words have been presented. The result is an ordered list that tends to change gradually along the selected feature dimension.

Adaptive condition

We designed the *adaptive* experimental condition to study the effect on memory for information that matched (or mismatched) the ways participants "naturally" organized their memories of the lists they studied. Like the other conditions, all participants in the adaptive condition studied a total of 16 lists, in a randomized order. We varied the words' colors and locations for every word presentation, as in the feature rich and order manipulation conditions.

All participants in the adaptive condition began the experiment by studying a set of four *initialization* lists. Words and features on these lists were presented in a randomized order (computed independently for each participant). These initialization lists were used to estimate each participant's "memory fingerprint," defined below. At a high level, a participant's memory fingerprint describes how they prioritize different semantic, lexico-

graphic, and/or visual features when they organize their memories.

Next, participants studied a sequence of 12 lists in three batches of 4 lists each. These batches came in three types: *random*, *stabilize*, and *destabilize*. The batch types determined how words on the lists in that batch were ordered. Lists in each batch were always presented consecutively (e.g., a participant might receive four random lists, followed by four stabilize lists, followed by four destabilize lists). The batch orders were evenly counterbalanced across participants: there are six possible orderings of the three batches, and 10 participants were randomly assigned to each ordering sub-condition.

Lists in the random batches were sorted randomly (as on the initialization lists and in the feature rich condition). Lists in the stabilize and destabilize batches were sorted in ways that either matched or mismatched each participant's memory fingerprint, respectively. Our procedures for computing participants' memory fingerprints and ordering the stabilize and destabilize lists are described next.

Feature clustering scores (uncorrected). Feature clustering scores describe participants' tendencies to recall similar presented items together in their recall sequences, where "similarity" considers one given feature dimension (e.g., category, color, etc.). We base our main approach to computing clustering scores on analogous temporal and semantic clustering scores developed by[?]. Computing the clustering score for one feature dimension starts by considering those feature values from the first word the participant recalled on the list. Next, we sort all not-yet-recalled words in ascending order according to their feature-based distance to the just-recalled item (see *Defining feature-based distances*). We then compute the percentile rank of the observed next recall. We averaged these percentile ranks across all of the participant's recalls for the current list to obtain a single uncorrected clustering score for the list, for the given feature dimension. We repeat this process for each feature dimension in turn to obtain a single uncorrected clustering score for each list,

for each feature dimension.

Permutation-corrected feature clustering scores. Suppose that two lists contain unequal numbers of items of each size. For example, suppose that list A contains all "large" items, whereas list B contains an equal mix of "large" and "small" items. For a participant recalling list A, any correctly recalled item will necessarily match the size of the previous correctly recalled item. In other words, successively recalling several list A items of the same size is essentially meaningless, since *any* correctly recalled list A word will be large. In contrast, successively recalling several list B items *could* be meaningful, since (early in the recall sequence) the yet-to-be-recalled items come from a mix of sizes. However, once all of the "small" items on list B have been recalled, the best possible next matching recall will be a large item. And all subsequent correct recalls must also be large items– so for those later recalls it becomes difficult to determine whether the participant is successively recalling "large" items because they are organizing their memories according to size, or (alternatively), whether they are simply recalling the yet-to-be-recalled items in a random order. In general, the precise order and blend of feature values expressed in a given list, the orders and numbers of correct recalls a participant makes, the number of itervening presentation positions between successive recalls, and so on, can all affect the range of clustering scores that are possible to observe for a given list. The uncorrected clustering score therefore conflates participants' actual memory organization with other "nuisance" factors.

Following our prior work[?], we used a permutation-correction procedure to help isolate the behavioral aspects of clustering that we were most interested in. After computing the uncorrected clustering score (for the given list and observed recall sequence), we compute a "null" distribution of n additional clustering scores after randomly shuffling the recall order (we use $n = 500$ in the present study). This null distribution represents an

approximation of the range of clustering scores one might expect to observe by "chance," given that a hypothetical participant was *not* truly clustering their recalls, but where the hypothetical participant studied and recalled exactly the same items (with the same features) as the true participant. We define the permutation-corrected clustering score as the percentile rank of the observed uncorrected clustering score in this estimated null distribution. In this way, a corrected score of 1 indicates that the observed score was greater than any clustering score one might expect by chance; in other words, good evidence that the participant was truly clustering their recalls along the given feature dimension.

Memory fingerprints.

Online "fingerprint" analysis.

Ordering "stabilize" lists by an estimated fingerprint.

Ordering "destabilize" lists by an estimated fingerprint.

Analyses

Probability of n^{th} recall curves

Probability of first recall curves[???] reflect the probability that an item will be recalled first, as a function of its serial position during encoding. To carry out this analysis, we initialized (for each participant) a number-of-lists (16) by number-of-words-per-list (16) matrix of zeros. Then, for each list, we found the index of the word that was recalled first, and we filled in that position in the matrix with a 1. Finally, we averaged over the rows of the matrix to obtain a 1 by 16 array of probabilities, for each participant. We used an analogous procedure to compute probabilility of n^{th} reacall curves for each participant.

Specifically, we filled in the corresponding matrices according to the n^{th} recall on each list that each participant made. When a given participant had made fewer than n recalls for a given list, we simply excluded that list from our analysis when computing that paritcipant's curve(s).

Lag-conditional response probability curve

The lag-conditional probability (lag-CRP) curve[?] reflects the probability of recalling a given item after the just-recalled item, as a function of their relative encoding positions (lag). In other words, a lag of 1 indicates that a recalled item was presented immediately after the previously recalled item, and a lag of 3 indicates that a recalled item came 3 items before the previously recalled item. For each recall transition (following the first recall), we computed the lag between the just-recalled word's presentation position and the next-recalled word's presentation position. We computed the proportions of transitions (between successively recalled words) for each lag, normalizing for the total numbers of possible transitions. In carrying out this analysis, we excluded all incorrect recalls and successive repetitions (e.g., recalling the same word twice in a row). This yielded, for each list, a 1 by number-of-lags (−15 to +15; 30 lags in total, excluding lags of 0) array of conditional probabilities. We averaged these probabilities across lists to obtain a single lag-CRP for each participant.

Temporal clustering score. Temporal clustering describes a participant's tendency to organize their recall sequences by the learned items' encoding positions. For instance, if a participant recalled the episode events in the exact order they occurred (or in exact reverse order), this would yield a score of 1. If a participant recalled the events in random order, this would yield an expected score of 0.5. For each recall-event transition (and separately for each participant), we sorted all not-yet-recalled events according to their

absolute lag (that is, distance away in the episode). We then computed the percentile rank of the next event the participant recalled. We took an average of these percentile ranks across all of the participant's recalls to obtain a single (uncorrected) temporal clustering score for the participant. Finally, as with the feature-based clustering scores, we applied our permutation correction procedure to obtain a corrected clustering score.

Computing low-dimensional embeddings of memory fingerprints

Serial position curve

Serial position curves[?] reflect the proportion of participants who remember each item as a function of the item's serial position during encoding. For each participant, we initialized a number-of-lists (16) by number-of-words-per-list (16) matrix of zeros. Then, for each correct recall, we identified the presentation position of the word and entered a 1 into that position (row: list; column: presentation position) in the matrix. This resulted in a matrix whose entries indicated whether or not the words presented at each position, on each list, were recalled by the participant (depending on whether the corresponding entires were set to one or zero). Finally, we averaged over the rows of the matrix to yield a 1 by 16 array representing the proportion of words at each position that the participant remembered.

Identifying event boundaries

Results

Figure S3.

Figure S7.

Figure S4.

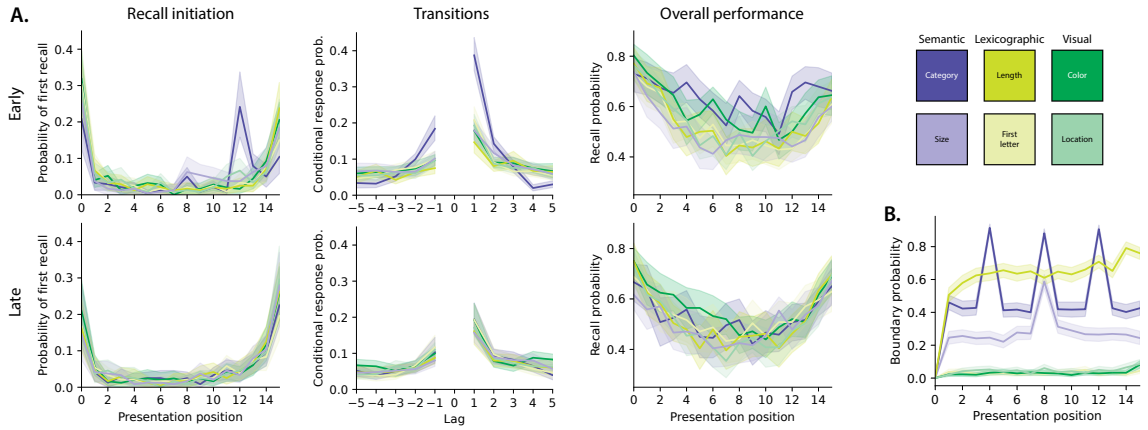

Figure 3: Recall dynamics in feature rich free recall (order manipulation conditions). A. Behavioral plots. **Left panels.** The probabilities of initiating recall with each word are plotted as a function of presentation position. **Middle panels.** The conditional probabilities of recalling each word are plotted as a function of the relative position (Lag) to the words recalled just-prior. **Right panels.** The overall probabilities of recalling each word are plotted as a function of presentation position. **All panels.** Error ribbons denote bootstrap-estimated 95% confidence intervals (calculated across participants). Top panels display the recall dynamics for early (order manipulation) lists in each condition (color). Bottom panels display the recall dynamics for late (randomly ordered) lists. See Figures S1 and S2 for analogous plots for the random (control) and adaptive conditions. **B.** Proportion of event boundaries (see *Methods*) for each condition's feature of focus, plotted as a function of presentation position.

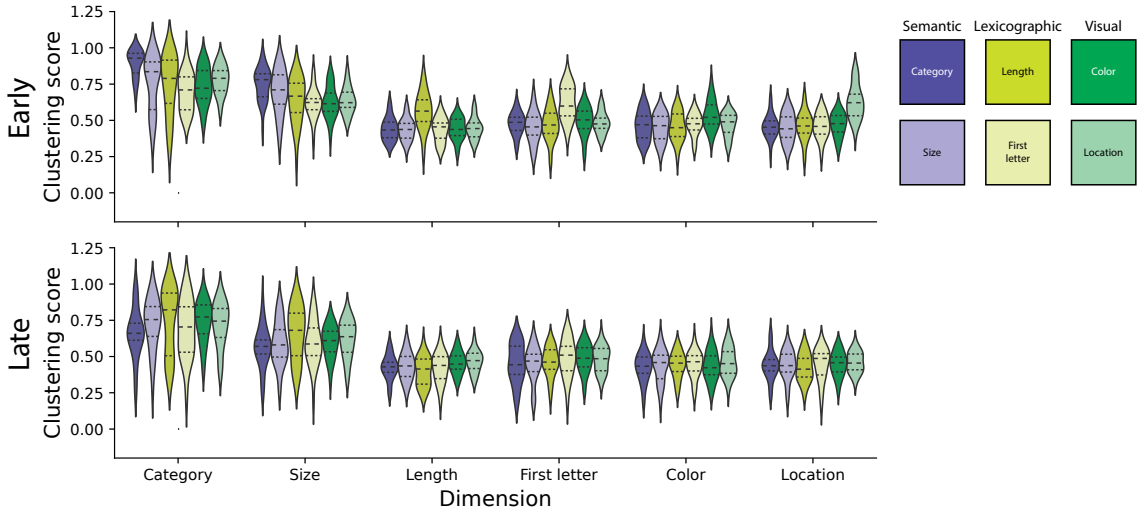

Figure 4: Memory "fingerprints" (order manipulation conditions). The across-participant distributions of clustering scores for each feature type (x-coordinate) are displayed for each experimental condition (color), separately for order manipulation (early, top) and randomly ordered (late, bottom) lists. See Figures S5 and S6 for analogous plots for the random (control) and adaptive conditions.

Discussion

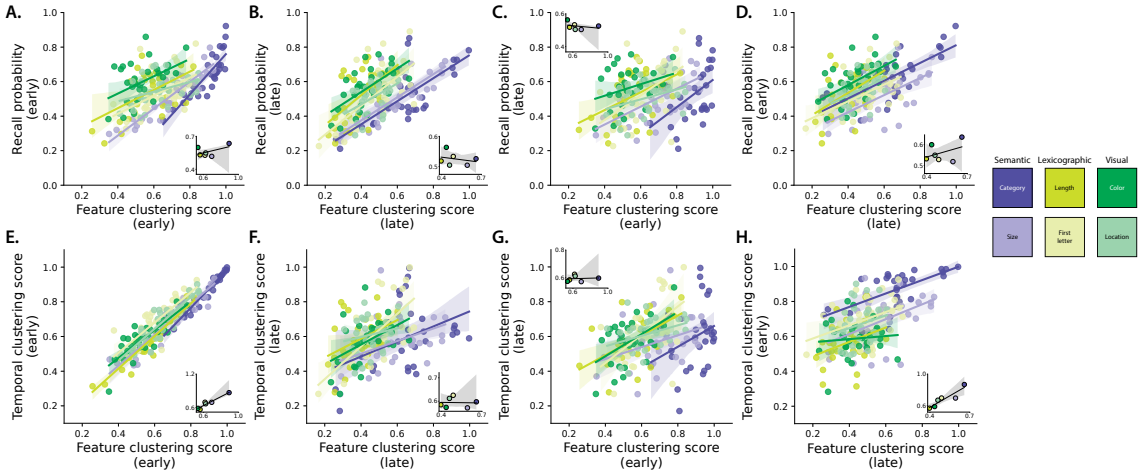

Figure 5: Interactions between feature clustering, recall probability, and contiguity. A. Recall probability versus feature clustering scores for order manipulation (early) lists. **B.** Recall probability versus feature clustering for randomly ordered (late) lists. **C.** Recall probability on late lists versus feature clustering on early lists. **D.** Recall probability on early lists versus feature clustering on late lists. **E.** Temporal clustering scores (contiguity) versus feature clustering scores on early lists. **F.** Temporal clustering scores versus feature clustering scores on late lists. **G.** Temporal clustering scores on late lists versus feature clustering scores on early lists. **H.** Temporal clustering scores on early lists versus feature clustering scores on late lists. **All panels.** Each dot in the main scatterplots denotes the average scores for one participant. The colored regression lines are computed across participants. The inset displays condition-averaged results, where each dot reflects a single condition and the regression line is computed across experimental conditions. All error ribbons denote bootstrap-estimated 95% confidence intervals.

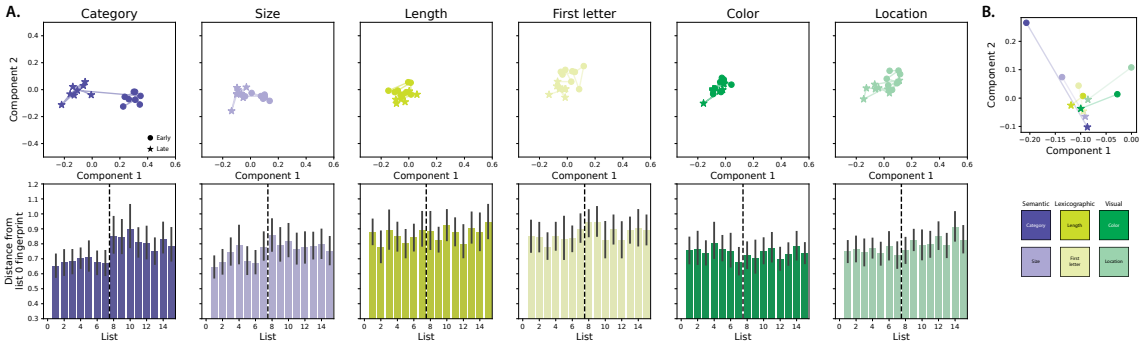

Figure 6: Memory fingerprint dynamics (order manipulation conditions). A. Each column (and color) reflects an experimental condition. In the top panels, each marker displays a 2D projection of the (across-participant) average memory fingerprint for one list. Order manipulation (early) lists are denoted by circles and randomly ordered (late) lists are denoted by stars. All of the fingerprints (across all conditions and lists) are projected into a common space. The bar plots in the bottom panels display the Euclidean distances of the per-list memory fingerprints to the list 0 fingerprint, for each condition. Error bars denote bootstrap-estimated 95% confidence intervals. The dotted vertical lines denote the boundaries between early and late lists. **B.** In this panel, the fingerprints for early (circle) and late (star) lists are averaged across lists and participants before projecting the fingerprints into a (new) 2D space. See Figure S8 for analogous plots for the random (control) conditions.

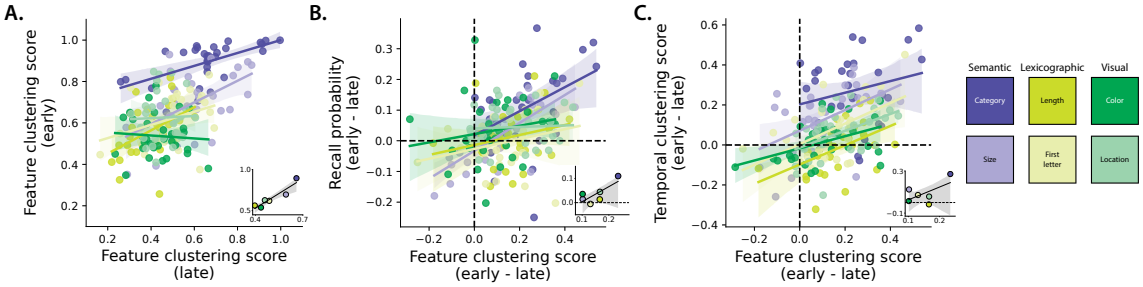

Figure 7: Feature clustering carryover effects. A. Feature clustering scores for ordder manipulation (early) versus randomly ordered (late) lists. **B.** Accuracy differences (on early versus late lists) versus feature clustering "carryover" (defined as the differences between the average clustering scores on early and late lists). **C.** Temporal clustering differences (on early versus late lists) versus feature clustering carryover. **All panels.** Each dot in the main scatterplots denotes the average scores for one participant. The colored regression lines are computed across participants. The inset displays condition-averaged results, where each dot reflects a single condition and the regression line is computed across experimental conditions. All error ribbons denote bootstrap-estimated 95% confidence intervals.

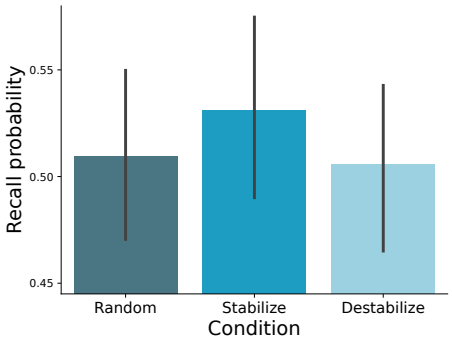

Figure 8: Recall performance (adaptive conditons). The bars display the average probability of recall (taken across words, lists, and participants) for lists from each adaptive condition. Error bars denote bootstrap-estimated 95% confidence intervals. For additional details about participants' behavior and performance during the adaptive conditions, see Figure S2.