

Carryover effects in free recall reveal how prior experiences influence memories of new experiences

Jeremy R. Manning^{1,*}, Kirsten Ziman^{1,2}, Emily Whitaker¹,
Paxton C. Fitzpatrick¹, Madeline R. Lee¹, Allison M Frantz¹,
Bryan J. Bollinger¹, Campbell E. Field¹, and Andrew C. Heusser^{1,3}

¹Dartmouth College

²Princeton University

³Akili Interactive

*Corresponding author: jeremy.r.manning@dartmouth.edu

Abstract

We perceive, interpret, and remember ongoing experiences through the lens of our prior experiences. Inferring that we are one type of situation versus another can lead us to interpret the same physical experience differently. In turn, this can affect how we focus our attention, form expectations of what will happen next, remember what is happening now, draw on our prior related experiences, and so on. To study these phenomena, we asked participants to perform simple word list learning tasks. Across different experimental conditions, we held the set of to-be-learned words constant, but we manipulated the orders in which the words were studied. We found that these order manipulations affected not only how the participants recalled the ordered lists, but also how they recalled later randomly ordered lists. Our work shows how structure in our ongoing experiences can exert influence on how we remember unrelated subsequent experiences.

17 Introduction

18 Experience is subjective: different people who encounter identical physical experiences
19 can take away very different meanings and memories. One reason is that our subjective
20 experiences in the moment are shaped in part the idiosyncratic prior experiences, mem-
21 ories, goals, thoughts, expectations, and emotions that we bring with us into the present
22 moment. These factors collectively define a *context* for our experiences¹⁹.

23 The contexts we encounter help us to construct *situation models*^{22,34} or *schemas*^{3,25} that
24 describe how experiences are likely to unfold based on our prior experiences with similar
25 contextual cues. For example, when we enter a sit-down restaurant, we might expect
26 to be seated at a table, given a menu, and served food. Priming someone to expect a
27 particular situation or context can also influence how they resolve potential ambiguities
28 in their ongoing experiences, including ambiguous movies and narratives⁴⁵.

29 Our understanding of how we form situation models and schemas, and how they in-
30 teract with our subjective experiences and memories, is constrained in part by substantial
31 differences in how we study these processes. Situation models and schemas are most often
32 studied using “naturalistic” stimuli such as narratives and movies^{28,47,48}. In contrast, our
33 understanding of how we organize our memories has been most widely studied using
34 more traditional paradigms like free recall of random word lists¹⁷. In free recall, partici-
35 pants study lists of items and are instructed to recall the items in any order they choose.
36 The orders in which words come to mind can provide insights into how participants have
37 organized their memories of the studied words. Because random word lists are unstruc-
38 tured by design, it is not clear if or how non-trivial situation models might apply to these
39 stimuli. Nevertheless, there are *some* commonalities between memory for word lists and
40 memory for real-world experiences.

41 Like remembering real-world experiences, remembering words on a studied list re-

quires distinguishing the current list from the rest of one's experience. To model this fundamental memory capability, cognitive scientists have posited the existence of a special representation, called *context*, that is associated with each list. According to early theories e.g.^{1,8} context representations are composed of many features which fluctuate from moment to moment, slowly drifting through a multidimensional feature space. During recall, this representation forms part of the retrieval cue, enabling us to distinguish list items from non-list items. Understanding the role of context in memory processes is particularly important in self-cued memory tasks, such as *free recall*, where the retrieval cue is "context" itself.

Over the past half-century, context-based models have enjoyed impressive success at explaining many stereotyped behaviors observed during free recall and other list-learning tasks^{8,10,14,18,29,30,32,37-39}. These phenomena include the well-known recency and primacy effects (superior recall of items from the end and, to a lesser extent, from the beginning of the study list), as well as semantic and temporal clustering effects⁷. The contiguity effect is an example of temporal clustering, which is perhaps the dominant form of organization in free recall. This effect can be seen in the tendency for people to successively recall items that occupied neighboring positions in the study list. For example, if a list contained the sub-sequence "ABSENCE HOLLOW PUPIL" and the participant recalls the word "HOLLOW", it is far more likely that the next response will be either "PUPIL" or "ABSENCE" than some other list item¹⁶. In addition, there is a strong forward bias in the contiguity effect: subjects make forward transitions (i.e., "HOLLOW" followed by "PUPIL") about twice as often as they make backward transitions, despite an overall tendency to begin recall at the end of the list. There are also striking effects of semantic clustering^{4,5,15,21,35}, whereby the recall of a given item is more likely to be followed by recall of a similar or related item than a dissimilar or unrelated one. In general, people organize memories for words along a

67 wide variety of stimulus dimensions. As captured by models like the *Context Maintenance*
68 *and Retrieval Model*³⁰, the stimulus features associated with each word (e.g. the word's
69 meaning, font size, font color, location on the screen, size of the object the word represents,
70 etc.) are incorporated into the participant's mental context representation^{19,22-24,40}. During
71 a memory test, any of these features may serve as a memory cue, which in turn leads the
72 participant to recall in succession words that share stimulus features.

73 A key mystery is whether the sorts of situation models and schemas that people use to
74 organize their memories of real-world experiences might map onto the clustering effects
75 that reflect how people organize their memories for word lists. On one hand, situation
76 models and clustering effects both reflect statistical regularities in ongoing experience.
77 Our memory systems exploit these regularities when generating inferences about the
78 unobserved past and yet-to-be-experienced future^{6,26,34,36,44}. On the other hand, the rich
79 structure of real-world experiences and other naturalistic stimuli that enable people to
80 form deep and meaningful situation models and schemas have no obvious analog in
81 simple word lists. Often lists in free recall studies are explicitly *designed* to be devoid of
82 exploitable temporal structure, for example by sorting the words in a random order¹⁷.

83 We designed an experimental paradigm to explore how people organize their mem-
84 ories for simple stimuli (word lists) whose temporal properties change across different
85 "situations," analogous to how the content of real-world experiences change across dif-
86 ferent real-world situations. We asked participants to study and freely recall a series
87 of word lists (Fig. 1). Across the different conditions in the experiment, we varied the
88 lists' presentation orders in different ways across lists. The studied items (words) were
89 designed to vary along three general dimensions: semantic (word *category*, and physical
90 *size of the referent*), lexicographic (word *length* and *first letter*), and visual (font *color* and
91 the onscreen *location* of each word). In our main manipulation conditions, we asked par-

102 participants to study and recall eight lists whose items were sorted by a target feature (e.g.,
103 word category). Next, we asked them to study and recall an additional eight lists whose
104 items had the same features, but that were sorted in a random temporal order. We were in-
105 terested in how these order manipulations affected participants' recall behaviors on early
106 (sorted) lists, as well as how order manipulations on early lists affected recall behaviors
107 on later (unsorted) lists. We used a series of control conditions as a baseline; in these
108 control conditions all of the lists were sorted randomly, but we manipulated the presence
109 or absence of the visual features. Finally, in an *adaptive* experimental condition we used
110 participants' recall behaviors on early lists to manipulate, in real-time, the presentation
111 orders of subsequent lists. In this adaptive condition, we sought to identify potential
112 commonalities within and across participants in how people organized their memories
113 and how those organizational tendencies affect overall performance.

104 **Materials and methods**

105 **Participants**

106 We enrolled a total of 491 Dartmouth undergraduate students across 11 experimental
107 conditions. The conditions included two primary controls (feature rich, reduced), two
108 secondary controls (reduced (early), reduced (late)), six order manipulation conditions
109 (category, size, length, first letter, color, and location), and a final adaptive condition. Each
110 of these conditions are described in the *Experimental design* subsection below.

111 Participants received course credit for enrolling in our study. We asked each participant
112 to fill out a demographic survey that included information about their self-reported age,
113 gender, ethnicity, race, education, vision, reading impairments, medications or recent
114 injuries, coffee consumption on the day of testing, and level of alertness at the time of

115 testing. All components of the demographics survey were optional. One participant
116 elected not to fill out any part of the demographic survey, and all other participants report
117 some or all of their requested demographic information.

118 We aimed to run (to completion) at least 60 participants in each of the two primary
119 control conditions and in the adaptive condition. In all other conditions we set a target
120 enrollment of at least 30 participants. Because our data collection efforts were coordinated
121 12 researchers and multiple testing rooms and computers, it was not feasible for individ-
122 ual experimenters to know how many participants had been run in each experimental
123 condition until the relevant databases were synchronized at the end of each working day.
124 We also over-enrolled participants for each condition to help ensure that we met our min-
125 imum enrollment targets even if some participants dropped out of the study prematurely
126 or did not show up for their testing session. This led us to exceed our target enrollments
127 for several conditions.

128 Participants were assigned to experimental conditions based loosely on their date of
129 participation. (This aspect of our procedure helped us to more easily synchronize the ex-
130 periment databases across multiple testing computers.) Of the 490 participants who opted
131 to fill out the demographics survey, reported ages ranged from 17 to 31 years (mean: 19.1;
132 standard deviation: 1.356). A total of 318 participants reported their gender as female,
133 170 as male, and 2 participants declined to report their gender. A total of 442 participants
134 reported their ethnicity as “not Hispanic or Latino,” 39 as “Hispanic or Latino,” and 9
135 declined to report their ethnicity. Participants reported their races as White (345 partic-
136 ipants), Asian (120 participants), Black or African American (31 participants), American
137 Indian or Alaska Native (11 participants), Native Hawaiian or Other Pacific Islander (4
138 participants), Mixed race (3 participants), Middle Eastern (1 participant), and Arab (1
139 participant). A total of 5 participants declined to report their race. We note that several

140 participants reported more than one of racial category. Participants reported their high-
141 est degrees achieved as “Some college” (359 participants), “High school graduate” (117
142 participants), “College graduate” (7 participants), “Some high school” (5 participants),
143 “Doctorate” (1 participant), and “Master’s degree” (1 participant). A total of 482 partici-
144 pants reported no reading impairments, and 8 reported mild reading impairments such
145 as mild dyslexia. A total of 489 participants reported having normal color vision and 1
146 participant reported that they were color blind. A total of 482 participants reported taking
147 no prescription medications and having no recent injuries; 4 participants reported having
148 ADHD, 1 reported having dyslexia, 1 reported having allergies, 1 reported a recently
149 torn ACL/MCL, and 1 reported a concussion from several months prior. The participants
150 reported consuming 0 – 3 cups of coffee prior to the testing session (mean: 0.32 cups;
151 standard deviation: 0.58 cups). Participants reported their current level of alertness, and
152 we converted their responses to numerical scores as follows: “very sluggish” (-2), “a little
153 sluggish” (-1), “neutral” (0), “a little alert” (1), and “very alert” (2). Across all partici-
154 pants, the full range of alertness levels were reported (range: -2 – 2; mean: 0.35; standard
155 deviation: 0.89).

156 We dropped from our dataset the 1 participant who reported abnormal color vision, as
157 well as 39 participants whose data were corrupted due to technical failures while running
158 the experiment or during the daily database merges. In total, this left usable data from
159 452 participants, broken down by experimental condition as follows: feature rich (67
160 participants), reduced (61 participants), reduced (late) (41 participants), reduced (early),
161 (42 participants), category (30 participants), size (30 participants), length (30 participants),
162 first letter (30 participants), color (31 participants), location (30 participants), and adaptive
163 (60 participants). The participant who declined to fill out their demographic survey
164 participated in the location condition, and we verified verbally that they had normal color

165 vision.

166 **Experimental design**

167 Our experiment is a variant of the classic free recall paradigm that we term *feature-rich free*
168 *recall*. In feature-rich free recall, participants study 16 lists, each comprised of 16 words that
169 vary along a number of stimulus dimensions (Fig. 1). The stimulus dimensions include
170 two semantic features related to the *meanings* of the words (semantic category, referent
171 object size), two lexicographic features related to the *letters* that make up the words (word
172 length in number of letters, identity of the word’s first letter), and two visual features
173 that are independent of the words themselves (text color, presentation location). Each
174 list contains four words from each of four different semantic categories and two object
175 sizes; all other stimulus features are randomized. After studying each list, the participant
176 attempts to recall as many words as they can from that list, in any order they choose.
177 Because each individual word is associated with several well-defined (and quantifiable)
178 features, and because each list incorporates a diverse mix of feature values along each
179 dimension, this allows us to evaluate participants’ memory fingerprints in rich detail.

180 **Stimuli**

181 Stimuli in our paradigm were 256 English words selected in a previous study⁴⁶. The words
182 all referred to concrete nouns, and were chosen from 15 unique semantic categories: body
183 parts, building-related, cities, clothing, countries, flowers, fruits, insects, instruments,
184 kitchen-related, mammals, (US) states, tools, trees, and vegetables. We also tagged each
185 word according to the approximate size of the object the word referred to. Words were
186 labeled as “small” if the corresponding object was likely able to “fit in a standard shoebox”
187 or “large” if the object was larger than a shoebox. Semantic categories varied in how many



Figure 1: Feature-rich free recall. After studying lists comprised of words that vary along several feature dimensions, participants verbally recall words in any order (microphone icon). Each experimental condition manipulates word features and/or presentation orders within and/or across lists. The rows display representative (illustrated) examples of the first lists participants might encounter in each condition. The rectangles during the “Presentation phase” show illustrated screen captures during a series of word presentations. Each word appeared onscreen for 2 seconds, followed by 2 seconds of blank screen. The red microphone icons during the “Recall” phase denote the one minute verbal recall interval. The labels on the right (and corresponding groupings on the left) denote experimental condition labels.

188 object sizes they reflected (mean number of different sizes per category: 1.33; standard
189 deviation: 0.49). The numbers of words in each semantic category also varied from 12
190 – 28 (mean number of words per category: 17.07; standard deviation number of words:
191 4.65). We also identified lexicographic features for each word, including the words' first
192 letters and lengths (i.e., number of letters). Across all categories, all possible first letters
193 were represented except for 'Q' (average number of unique first letters per category: 11;
194 standard deviation: 2 letters). Word lengths ranged from 3 – 12 letters (average: 6.17
195 letters; standard deviation: 2.06 letters).

196 We assigned the categorized words into a total of 16 lists with several constraints.
197 First, we required that each list contained words from exactly 4 unique categories, each
198 with exactly 4 exemplars from each category. Second, we required that (across all words
199 on the list) at least one instance of both object sizes were represented. On average, each
200 category was represented in 4.27 lists (standard deviation: 1.16 lists). Aside from these
201 two constraints, we assigned each word to a unique list. After random assignment, each
202 list contained words with an average of 11.13 unique starting letters (standard deviation:
203 1.15 letters) and an average word length of 6.17 letters (standard deviation: 0.34 letters).

204 The above assignments of words to lists was performed once across all participants,
205 such that every participant studied the same set of 16 lists. In every condition we random-
206 ized the study order of these lists across participants. For participants in some conditions,
207 on some lists, we also randomly varied two additional visual features to each word: the
208 presentation font color, and the word's onscreen location. These attributes were assigned
209 independently for word (and for every participant) at the times the words were displayed
210 onscreen. These visual features were varied for words in all lists and conditions except for
211 the "reduced" condition (all lists), the first eight lists of the "reduced (early)" condition,
212 and the last eight lists of the "reduced (late)" condition. In these latter cases, words were

213 all presented in black at the center of the experimental computer’s display.

214 To assign a random font color to each word, we selected three integers uniformly
215 and at random between 0 and 255, corresponding to the red (r), green (g), and blue (b)
216 color channels for that word. To assign random presentation locations to each word, we
217 selected two floating point numbers uniformly at random (one for the word’s horizontal
218 x coordinate and the other for its vertical y coordinate). The bounds of these coordinates
219 were selected to cover the entire visible area of the display without cutting off any part of
220 the words. The words were shown on 27 in (diagonal) Retina 5K iMac displays (resolution:
221 5120×2880 pixels).

222 Most of the experimental manipulations we carried out entailed presenting or sorting
223 the presented words differently on the first eight lists participants studied (which we call
224 *early* lists) versus on the final eight lists they studied (*late* lists). Since every participant
225 studied exactly 16 lists, using this terminology every list was either “early” or “late”
226 depending on its order in the list study sequence.

227 **Real-time speech-to-text processing**

228 Our experimental paradigm incorporates the Google Cloud Speech API speech-to-text
229 engine¹¹ to automatically transcribe participants’ verbal recalls into text. This allows
230 recalls to be transcribed in real time— a distinguishing feature of the experiment; in typical
231 verbal recall experiments the audio data must be parsed manually. In prior work, we
232 used a similar experimental setup (equivalent to the “reduced” condition in the present
233 study) to verify that the automatically transcribed recalls were sufficiently close to human-
234 transcribed recalls to yield reliable data⁴⁶. This real-time speech processing component of
235 the paradigm plays an important role in the “adaptive” condition of the experiment, as
236 described below.

237 **Random conditions (Fig. 1, top four rows)**

238 We used four “control” conditions to evaluate and explore participants’ baseline behaviors.
239 We also used performance on these control conditions to help interpret performance in
240 other “manipulation” conditions. Two control conditions served as “anchorpoints.” In the
241 first anchorpoint condition, which we call the *feature rich* condition, we randomly shuffled
242 the presentation order (independently for each participant) of the words on each list. In
243 the second anchorpoint condition, which we call the *reduced* condition, we randomized
244 word presentations as in the feature rich condition. However, rather than assigning each
245 word a random color and location, we instead displayed all of the words in black and at
246 the center of the screen.

247 In the *reduced (early)* condition, we followed the “reduced” procedure (presenting each
248 word in black at the center of the screen) for early lists, and followed the “feature rich”
249 procedure (presenting each word in a random color and location) for late lists. Finally, in
250 the *reduced (late)* condition, we followed the feature rich procedure for earlylists and the
251 reduced procedure for late lists.

252 **Order manipulation conditions (Fig. 1, middle six rows)**

253 Each of six *order manipulation* conditions used a different feature-based sorting procedure
254 to order words on early lists, where each sorting procedure relied on one relevant feature
255 dimension. All of the irrelevant features varied freely across words on early lists, in
256 that we did not consider irrelevant features in ordering the early lists. However, some
257 features were correlated– for example, some semantic categories of words referred to
258 objects that tended to be a particular size, which means that category and size are not
259 fully independent. On late lists, the words were always presented in a randomized order
260 (chosen anew for each participant). In all of the order manipulation conditions, we varied

261 words' font colors and onscreen locations, as in the feature rich condition.

262 **Defining feature-based distances.** Sorting words according to a given relevant feature
263 requires first defining a distance function for quantifying the dissimilarity between each
264 pair of features. This function varied according to the type of features. Semantic features
265 (category and size) are *categorical*. For these features, we defined a binary distance function:
266 two words were considered to “match” (i.e., have a distance of 0) if their labels are the
267 same (i.e., both from the same semantic category or both of the same size). If two words'
268 labels were different for a given feature, we defined the words to have a distance of 1 for
269 that feature. Lexicographic features (length and first letter) are *discrete*. For these features
270 we defined a discrete distance function. Specifically, we defined the distance between
271 two words as either the absolute difference between their lengths, or the absolute distance
272 between their starting letters in the English alphabet, respectively. For example, two
273 words that started with the same letter would have a “first letter” distance of 0, and words
274 starting with ‘J’ and ‘A’ would have a first letter distance of 9. Because words' lengths
275 and letters' positions in the alphabet are always integers, these discrete distances always
276 take on integer values. Finally, the visual features (color and location) are *continuous* and
277 *multivariate*, in that each “feature” takes on multiple (positive) real values. We defined the
278 “color” and “location” distances between two words as the Euclidean distances between
279 their (r, g, b) color or (x, y) location vectors, respectively. Therefore the color and location
280 distance measures always take on positive real values (upper bounded at 441.67 for color, or
281 27 in for location, reflecting the distances between the corresponding maximally different
282 vectors).

283 **Constructing feature-sorted lists.** Given a list of words, a relevant feature, and each
284 word's value(s) for that feature, we developed a stochastic algorithm for (noisily) sorting




Figure 2: Generating stochastic feature-sorted lists. For a given feature dimension (e.g., color), we compute the similarity (Eqn. 1) between the feature value(s) of the previous item, x , and all yet-to-be-presented items ($a - g$). Next, we normalize these similarity scores so that they sum to one. We lay in sequence a set of “sticks,” one for each candidate item, whose lengths are equal to these normalized similarity scores. Note that the combined lengths of these sticks is one. To select the next to-be-presented item, we draw a random number, r , from the uniform distribution bounded between 0 and 1 (inclusive). The identity of the next item is given by the stick adjacent to an indicator that moves distance r (starting from 0) along the sequence of sticks. In this case, the next to-be-presented item is e . Note that each item’s chances of selection is proportional to its similarity to the previous item, along the given feature dimension.

285 the words. First, we choose a word uniformly at random from the set of candidates. Next,
 286 we compute the distances between the chosen word’s feature(s) and the corresponding
 287 feature(s) of all yet-to-be-presented words. Third, we convert these distances (between the
 288 previously presented word’s feature values, a , and the candidate word’s feature values, b)
 289 to similarity scores:

$$\text{similarity}(a, b) = \exp\{-\tau \cdot \text{distance}(a, b)\}, \quad (1)$$

290 where $\tau = 1$ in our implementation. We note that increasing the value of τ would amplify
 291 the influence of similarity on order, and decreasing the value of τ would diminish the
 292 influence of similarity on order. Also note that this approach requires $\tau > 0$. Finally, we
 293 computed a set of normalized similarity values by dividing the similarities by their sum:

$$\text{similarity}_{\text{normalized}}(a, b) = \frac{\text{similarity}(a, b)}{\sum_{i=1}^n \text{similarity}(a, i)}, \quad (2)$$

294 where in the demoniator, i takes on each of the n feature values of the to-be-presented
 295 words. The resulting set of normalized similarity scores sums to one.

296 As illustrated in Figure 2, we use these normalized similarity scores to construct a

sequence of “sticks” that we lay end to end in a line. Each of the n sticks corresponds to a single to-be-presented word, and the stick lengths are proportional to the relative similarities between each word’s feature value(s) and the feature value(s) of the just-presented word. We choose the next to-be-presented word by moving an indicator along the set of sticks, by a distance chosen uniformly at random on the interval $[0,1]$. We select the word associated with the stick lying next to the indicator to be presented next. This process continues iteratively (re-computing the similarity scores and stochastically choosing the next to-be-presented word using the just-presented word) until all of the words have been presented. The result is an ordered list that tends to change gradually along the selected feature dimension.

Adaptive condition

We designed the *adaptive* experimental condition to study the effect on memory for information that matched (or mismatched) the ways participants “naturally” organized their memories of the lists they studied. Like the other conditions, all participants in the adaptive condition studied a total of 16 lists, in a randomized order. We varied the words’ colors and locations for every word presentation, as in the feature rich and order manipulation conditions.

All participants in the adaptive condition began the experiment by studying a set of four *initialization* lists. Words and features on these lists were presented in a randomized order (computed independently for each participant). These initialization lists were used to estimate each participant’s “memory fingerprint,” defined below. At a high level, a participant’s memory fingerprint describes how they prioritize different semantic, lexicographic, and/or visual features when they organize their memories.

Next, participants studied a sequence of 12 lists in three batches of 4 lists each. These

321 batches came in three types: *random*, *stabilize*, and *destabilize*. The batch types determined
322 how words on the lists in that batch were ordered. Lists in each batch were always
323 presented consecutively (e.g., a participant might receive four random lists, followed
324 by four stabilize lists, followed by four destabilize lists). The batch orders were evenly
325 counterbalanced across participants: there are six possible orderings of the three batches,
326 and 10 participants were randomly assigned to each ordering sub-condition.

327 Lists in the random batches were sorted randomly (as on the initialization lists and in
328 the feature rich condition). Lists in the stabilize and destabilize batches were sorted in
329 ways that either matched or mismatched each participant’s memory fingerprint, respec-
330 tively. Our procedures for computing participants’ memory fingerprints and ordering the
331 stabilize and destabilize lists are described next.

332 **Feature clustering scores (uncorrected).** Feature clustering scores describe participants’
333 tendencies to recall similar presented items together in their recall sequences, where “sim-
334 ilarity” considers one given feature dimension (e.g., category, color, etc.). We base our
335 main approach to computing clustering scores on analogous temporal and semantic clus-
336 tering scores developed by³⁰. Computing the clustering score for one feature dimension
337 starts by considering those feature values from the first word the participant recalled on
338 the list. Next, we sort all not-yet-recalled words in ascending order according to their
339 feature-based distance to the just-recalled item (see *Defining feature-based distances*). We
340 then compute the percentile rank of the observed next recall. We averaged these percentile
341 ranks across all of the participant’s recalls for the current list to obtain a single uncorrected
342 clustering score for the list, for the given feature dimension. We repeat this process for
343 each feature dimension in turn to obtain a single uncorrected clustering score for each list,
344 for each feature dimension.

345 **Temporal clustering score (uncorrected).** Temporal clustering describes a participant's
346 tendency to organize their recall sequences by the learned items' encoding positions. For
347 instance, if a participant recalled the episode events in the exact order they occurred (or
348 in exact reverse order), this would yield a score of 1. If a participant recalled the events in
349 random order, this would yield an expected score of 0.5. For each recall-event transition
350 (and separately for each participant), we sorted all not-yet-recalled events according to
351 their absolute lag (that is, distance away in the episode). We then computed the percentile
352 rank of the next event the participant recalled. We took an average of these percentile ranks
353 across all of the participant's recalls to obtain a single (uncorrected) temporal clustering
354 score for the participant.

355 **Permutation-corrected feature clustering scores.** Suppose that two lists contain unequal
356 numbers of items of each size. For example, suppose that list *A* contains all "large" items,
357 whereas list *B* contains an equal mix of "large" and "small" items. For a participant
358 recalling list *A*, any correctly recalled item will necessarily match the size of the previous
359 correctly recalled item. In other words, successively recalling several list *A* items of the
360 same size is essentially meaningless, since *any* correctly recalled list *A* word will be large.
361 In contrast, successively recalling several list *B* items *could* be meaningful, since (early in
362 the recall sequence) the yet-to-be-recalled items come from a mix of sizes. However, once
363 all of the "small" items on list *B* have been recalled, the best possible next matching recall
364 will be a large item. And all subsequent correct recalls must also be large items– so for
365 those later recalls it becomes difficult to determine whether the participant is successively
366 recalling "large" items because they are organizing their memories according to size, or
367 (alternatively), whether they are simply recalling the yet-to-be-recalled items in a random
368 order. In general, the precise order and blend of feature values expressed in a given list,
369 the orders and numbers of correct recalls a participant makes, the number of intervening

370 presentation positions between successive recalls, and so on, can all affect the range of
371 clustering scores that are possible to observe for a given list. The uncorrected clustering
372 score therefore conflates participants' actual memory organization with other "nuisance"
373 factors.

374 Following our prior work¹², we used a permutation-correction procedure to help iso-
375 late the behavioral aspects of clustering that we were most interested in. After computing
376 the uncorrected clustering score (for the given list and observed recall sequence), we com-
377 pute a "null" distribution of n additional clustering scores after randomly shuffling the
378 recall order (we use $n = 500$ in the present study). This null distribution represents an
379 approximation of the range of clustering scores one might expect to observe by "chance,"
380 given that a hypothetical participant was *not* truly clustering their recalls, but where the
381 hypothetical participant studied and recalled exactly the same items (with the same fea-
382 tures) as the true participant. We define the permutation-corrected clustering score as the
383 percentile rank of the observed uncorrected clustering score in this estimated null distri-
384 bution. In this way, a corrected score of 1 indicates that the observed score was greater
385 than any clustering score one might expect by chance; in other words, good evidence that
386 the participant was truly clustering their recalls along the given feature dimension. We
387 applied this correction procedure to all of the clustering scores (feature and temporal)
388 reported in this paper.

389 **Memory fingerprints.** We define each participant's *memory fingerprint* as the set of their
390 permutation-corrected clustering scores across all dimensions we tracked in our study,
391 including their six feature-based clustering scores (category, size, length, first letter, color,
392 and location) and their temporal clustering score. Conceptually, this memory fingerprint
393 describes the participant's tendencies to order (and, presumably, organize in memory)
394 the studied words along each dimension. To obtain stable estimates of these fingerprints

395 for each participant, we averaged clustering scores across lists. We also tracked and
396 characterized how participants' fingerprints changed across lists (e.g., Figs. 6, S8).

397 **Online “fingerprint” analysis.** The presentation orders of some lists in the adaptive
398 condition of our experiment (see *Adaptive condition*) were sorted according to participants'
399 *current* memory fingerprint, estimated using all of the lists they had studied up to that point
400 in the experiment. Because our experiment incorporated a speech-to-text component, all
401 of the behavioral data for each participant could be analyzed just a few seconds after
402 the conclusion of the recall intervals for each list. We used the Quail Python package¹²
403 to apply speech-to-text algorithms to the just collected data, aggregate the data for the
404 given participant, and estimate the participant's memory fingerprint using all of their
405 available data up to that point in the experiment. Two aspects of our implementation are
406 worth noting. First, because memory fingerprints are averaged across lists, the already-
407 computed memory fingerprints for earlier lists could be cached and loaded as needed
408 in future computations. This meant that our computations pertaining to updating our
409 estimates of a participant's memory fingerprint only needed to consider data from the
410 most recent list. Second, each element of the null distributions of uncorrected fingerprint
411 scores (see *Permutation-corrected feature clustering scores*) could be estimated independently
412 from the others. This enabled us to make use of the testing computers' multi-core CPU
413 architectures by elements of the null distributions in batches of eight (i.e., the number
414 of CPU cores on each testing computer). Taken together, we were able to compress
415 the fingerprint computations into just a few seconds of computing time. The combined
416 processing time for the speech-to-text algorithm and fingerprint computations easily fit
417 within the inter-list intervals, where participants typically paused before moving on to the
418 next list.

419 **Ordering “stabilize” and “destabilize” lists by an estimated fingerprint.** In the adap-
420 tive condition of our experiment, the presentation orders for *stabilize* and *destabilize* lists
421 were chosen to either maximally or minimally (respectively) comport with participants’
422 memory fingerprints. Given a participant’s memory fingerprint and a to-be-presented set
423 of items, we designed a permutation-based procedure for ordering the items. First, we
424 dropped from the participant’s fingerprint the temporal clustering score. For the remain-
425 ing feature dimensions, we arranged the clustering scores in the fingerprint into a template
426 vector, f . Second, we computed $n = 2500$ random permutations of the to-be-presented
427 items. These permutations served as prospective presentation orders. We sought to select
428 the specific order that most (or least) matched f . Third, for each random permutation, we
429 computed the (permutation-corrected) “fingerprint,” treating the permutation as though
430 it were a potential “perfect” recall sequence. (We did not include temporal clustering
431 scores in these fingerprints.) This yielded a “simulated fingerprint” vector, \hat{f}_p for each per-
432 mutation p . We used these simulated fingerprints to select a specific permutation, i , that
433 either maximized (for stabilized lists) or minimized (for destabilize lists) the correlation
434 between \hat{f}_i and f .

435 **Computing low-dimensional embeddings of memory fingerprints**

436 Following some of our prior work¹³, we use low-dimensional embeddings to help vi-
437 sualize how participants’ memory fingerprints change across lists (Figs. 6A, S8A). To
438 compute a shared embedding space across participants and experimental conditions, we
439 concatenated the full set of fingerprints (across all lists, participants, and experimental
440 conditions) to create a large matrix with number-of-lists \times number-of-participants rows
441 and seven columns (one for each word feature dimension’s clustering scores, plus an
442 additional temporal clustering score column). We used principal components analysis to

443 project the seven-dimensional observations into a two-dimensional space (using the two
444 principal components that explained the most variance in the data). For two visualizations
445 (Figs. 6B, and S8B) we computed an additional set of two-dimensional embeddings for par-
446 ticipants' *average* fingerprints (i.e., across lists within a given group of lists– early or late).
447 For those visualizations we averaged each participant's rows (for the given group of lists)
448 in the combined fingerprint matrix prior to projecting it into the shared two-dimensional
449 space. This yielded a single two-dimensional coordinate for each *participant* and *list group*,
450 rather than for each individual list. We used these embeddings solely for visualization.
451 All statistical tests were carried out in the original (seven-dimensional) feature spaces.

452 **Analyses**

453 **Probability of n^{th} recall curves**

454 Probability of first recall curves^{2,31,43} reflect the probability that an item will be recalled
455 first, as a function of its serial position during encoding. To carry out this analysis, we
456 initialized (for each participant) a number-of-lists (16) by number-of-words-per-list (16)
457 matrix of zeros. Then, for each list, we found the index of the word that was recalled first,
458 and we filled in that position in the matrix with a 1. Finally, we averaged over the rows
459 of the matrix to obtain a 1 by 16 array of probabilities, for each participant. We used an
460 analogous procedure to compute probability of n^{th} recall curves for each participant.
461 Specifically, we filled in the corresponding matrices according to the n^{th} recall on each
462 list that each participant made. When a given participant had made fewer than n recalls
463 for a given list, we simply excluded that list from our analysis when computing that
464 participant's curve(s).

465 **Lag-conditional response probability curve**

466 The lag-conditional probability (lag-CRP) curve¹⁶ reflects the probability of recalling a
467 given item after the just-recalled item, as a function of their relative encoding positions
468 (lag). In other words, a lag of 1 indicates that a recalled item was presented immediately
469 after the previously recalled item, and a lag of 3 indicates that a recalled item came 3
470 items before the previously recalled item. For each recall transition (following the first
471 recall), we computed the lag between the just-recalled word's presentation position and the
472 next-recalled word's presentation position. We computed the proportions of transitions
473 (between successively recalled words) for each lag, normalizing for the total numbers of
474 possible transitions. In carrying out this analysis, we excluded all incorrect recalls and
475 successive repetitions (e.g., recalling the same word twice in a row). This yielded, for
476 each list, a 1 by number-of-lags (-15 to +15; 30 lags in total, excluding lags of 0) array of
477 conditional probabilities. We averaged these probabilities across lists to obtain a single
478 lag-CRP for each participant.

479 **Serial position curve**

480 Serial position curves²⁷ reflect the proportion of participants who remember each item as a
481 function of the item's serial position during encoding. For each participant, we initialized
482 a number-of-lists (16) by number-of-words-per-list (16) matrix of zeros. Then, for each
483 correct recall, we identified the presentation position of the word and entered a 1 into that
484 position (row: list; column: presentation position) in the matrix. This resulted in a matrix
485 whose entries indicated whether or not the words presented at each position, on each list,
486 were recalled by the participant (depending on whether the corresponding entries were
487 set to one or zero). Finally, we averaged over the rows of the matrix to yield a 1 by 16 array
488 representing the proportion of words at each position that the participant remembered.

489 Identifying event boundaries

490 We used the distances between feature values for successively presented words (see *Defin-*
491 *ing feature-based distances*) to estimate “event boundaries” where the feature values changed
492 more than usual^{7,9,20,33,41,42}. For each list, for each feature dimension, we computed the
493 distribution of distances between the feature values for successively presented words. We
494 defined event boundaries (e.g., Fig. 3B) as occurring between any successive pair of words
495 whose distances along the given feature dimension were greater than one standard devia-
496 tion above the mean for that list. Note that, because event boundaries are defined for each
497 feature dimension, each individual list may contain several sets of event boundaries, each
498 at different moments in the presentation sequence (depending on the feature dimension
499 of interest).

500 Results

501 We sought to manipulate two aspects of how participants memorized sequences of word
502 lists. First, we added two additional sources of visual variation to the individual word
503 presentations: font color and onscreen location. Importantly, these visual features were
504 independent of the meaning or semantic content of the words (e.g., word category, size
505 of the referent) and of the lexicographic properties of the word (e.g., word length, first
506 letter). We wondered whether this additional word-independent information might facil-
507 itate recall (e.g., by providing new potential ways of organizing or retrieving memories
508 of the studied words) or impair recall (e.g., by distracting participants). Second, our pri-
509 mary experimental manipulations entailed manipulating the orders in which words were
510 studied (and how those orderings changed over time). We wondered whether presenting
511 the same list of words in different orders (e.g., sorted along one feature dimension versus

another) might serve to influence how participants organized their memories of the words. We also wondered whether some order manipulations might be temporally “sticky” by influencing how *future* lists were remembered.

- what do additional visual features add? (compare reduced vs. feature rich) - are the visual features “sticky”? (compare feature rich vs. reduced (early), also reduced vs. reduced (early)) - are impoverished stimuli “sticky”? (compare feature rich vs. reduced (late), reduced vs. reduced (late), also reduced (early) vs. reduced (late))

- are order effects “sticky”? compare behavior on early vs. late lists for order manipulation condition - does feature clustering on early lists correlate with recall on early (or late) lists? - does feature clustering on late lists correlate with recall on early (or late) lists? - (ditto, but replace “recall” with “temporal clustering”)

- for feature-rich lists, do order effects matter? - recall + dynamics + organization on order-manipulation conditions vs. feature rich - fingerprint trajectories: how much do fingerprints change over time, are they sensitive to order manipulations?

- are fingerprints malleable? how does match between fingerprint + presentation order affect recall performance (adaptive conditions)?

Figure S3.

Figure S7.

Figure S4.

Discussion

References

[1] Anderson, J. R. and Bower, G. H. (1972). Recognition and retrieval processes in free recall. *Psychological Review*, 79(2):97–123.




Figure 3: Recall dynamics in feature rich free recall (order manipulation conditions). **A.** Behavioral plots. **Left panels.** The probabilities of initiating recall with each word are plotted as a function of presentation position. **Middle panels.** The conditional probabilities of recalling each word are plotted as a function of the relative position (Lag) to the words recalled just-prior. **Right panels.** The overall probabilities of recalling each word are plotted as a function of presentation position. **All panels.** Error ribbons denote bootstrap-estimated 95% confidence intervals (calculated across participants). Top panels display the recall dynamics for early (order manipulation) lists in each condition (color). Bottom panels display the recall dynamics for late (randomly ordered) lists. See Figures S1 and S2 for analogous plots for the random (control) and adaptive conditions. **B.** Proportion of event boundaries (see *Methods*) for each condition's feature of focus, plotted as a function of presentation position.




Figure 4: Memory “fingerprints” (order manipulation conditions). The across-participant distributions of clustering scores for each feature type (x -coordinate) are displayed for each experimental condition (color), separately for order manipulation (early, top) and randomly ordered (late, bottom) lists. See Figures S5 and S6 for analogous plots for the random (control) and adaptive conditions.

- 535 [2] Atkinson, R. C. and Shiffrin, R. M. (1968). Human memory: A proposed system and its
536 control processes. In Spence, K. W. and Spence, J. T., editors, *The Psychology of Learning
537 and Motivation*, volume 2, pages 89–105. Academic Press, New York, NY.
- 538 [3] Baldassano, C., Hasson, U., and Norman, K. A. (2018). Representation of real-world
539 event schemas during narrative perception. *The Journal of Neuroscience*, 38(45):9689–
540 9699.
- 541 [4] Bousfield, W. A. (1953). The occurrence of clustering in the recall of randomly arranged
542 associates. *Journal of General Psychology*, 49:229–240.
- 543 [5] Bousfield, W. A., Sedgewick, C. H., and Cohen, B. H. (1954). Certain temporal charac-
544 teristics of the recall of verbal associates. *American Journal of Psychology*, 67:111–118.




Figure 5: Interactions between feature clustering, recall probability, and contiguity. A. Recall probability versus feature clustering scores for order manipulation (early) lists. B. Recall probability versus feature clustering for randomly ordered (late) lists. C. Recall probability on late lists versus feature clustering on early lists. D. Recall probability on early lists versus feature clustering on late lists. E. Temporal clustering scores (contiguity) versus feature clustering scores on early lists. F. Temporal clustering scores versus feature clustering scores on late lists. G. Temporal clustering scores on late lists versus feature clustering scores on early lists. H. Temporal clustering scores on early lists versus feature clustering scores on late lists. **All panels.** Each dot in the main scatterplots denotes the average scores for one participant. The colored regression lines are computed across participants. The inset displays condition-averaged results, where each dot reflects a single condition and the regression line is computed across experimental conditions. All error ribbons denote bootstrap-estimated 95% confidence intervals.

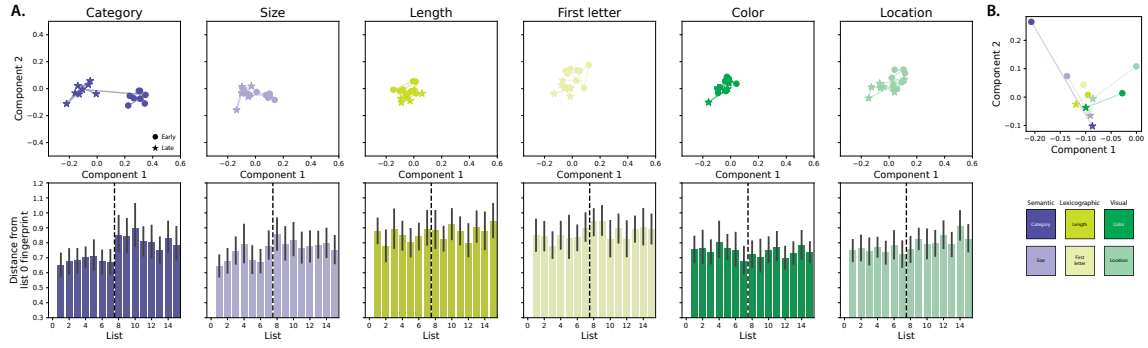


Figure 6: Memory fingerprint dynamics (order manipulation conditions). **A.** Each column (and color) reflects an experimental condition. In the top panels, each marker displays a 2D projection of the (across-participant) average memory fingerprint for one list. Order manipulation (early) lists are denoted by circles and randomly ordered (late) lists are denoted by stars. All of the fingerprints (across all conditions and lists) are projected into a common space. The bar plots in the bottom panels display the Euclidean distances of the per-list memory fingerprints to the list 0 fingerprint, for each condition. Error bars denote bootstrap-estimated 95% confidence intervals. The dotted vertical lines denote the boundaries between early and late lists. **B.** In this panel, the fingerprints for early (circle) and late (star) lists are averaged across lists and participants before projecting the fingerprints into a (new) 2D space. See Figure S8 for analogous plots for the random (control) conditions.




Figure 7: Feature clustering carryover effects. **A.** Feature clustering scores for order manipulation (early) versus randomly ordered (late) lists. **B.** Accuracy differences (on early versus late lists) versus feature clustering “carryover” (defined as the differences between the average clustering scores on early and late lists). **C.** Temporal clustering differences (on early versus late lists) versus feature clustering carryover. **All panels.** Each dot in the main scatterplots denotes the average scores for one participant. The colored regression lines are computed across participants. The inset displays condition-averaged results, where each dot reflects a single condition and the regression line is computed across experimental conditions. All error ribbons denote bootstrap-estimated 95% confidence intervals.




Figure 8: Recall performance (adaptive conditons). The bars display the average probability of recall (taken across words, lists, and participants) for lists from each adaptive condition. Error bars denote bootstrap-estimated 95% confidence intervals. For additional details about participants' behavior and performance during the adaptive conditions, see Figure S2.

- 545 [6] Bower, G. H., Black, J. B., and Turner, T. J. (1979). Scripts in memory for text. *Cognitive*
546 *Psychology*, 11(2):177–220.
- 547 [7] DuBrow, S. and Davachi, L. (2016). Temporal binding within and across events.
548 *Neurobiology of Learning and Memory*, 134:107–114.
- 549 [8] Estes, W. K. (1955). Statistical theory of spontaneous recovery and regression. *Psycho-*
550 *logical Review*, 62:145–154.
- 551 [9] Ezzyat, Y. and Davachi, L. (2011). What constitutes an episode in episodic memory?
552 *Psychological Science*, 22(2):243–252.
- 553 [10] Glenberg, A. M., Bradley, M. M., Kraus, T. A., and Renzaglia, G. J. (1983). Studies of
554 the long-term recency effect: support for a contextually guided retrieval theory. *Journal*
555 *of Experimental Psychology: Learning, Memory, and Cognition*, 12:413–418.
- 556 [11] Halpern, Y., Hall, K. B., Schogol, V., Riley, M., Roark, B., Skobeltsyn, G., and Bäuml,
557 M. (2016). Contextual prediction models for speech recognition. In *Interspeech*, pages
558 2338–2342.

- 559 [12] Heusser, A. C., Fitzpatrick, P. C., Field, C. E., Ziman, K., and Manning, J. R. (2017).
560 Quail: a Python toolbox for analyzing and plotting free recall data. *Journal of Open*
561 *Source Software*, 10.21105/joss.00424.
- 562 [13] Heusser, A. C., Ziman, K., Owen, L. L. W., and Manning, J. R. (2018). HyperTools:
563 a Python toolbox for gaining geometric insights into high-dimensional data. *Journal of*
564 *Machine Learning Research*, 18(152):1–6.
- 565 [14] Howard, M. W. and Kahana, M. J. (2002). A distributed representation of temporal
566 context. *Journal of Mathematical Psychology*, 46:269–299.
- 567 [15] Jenkins, J. J. and Russell, W. A. (1952). Associative clustering during recall. *Journal of*
568 *Abnormal and Social Psychology*, 47:818–821.
- 569 [16] Kahana, M. J. (1996). Associative retrieval processes in free recall. *Memory and*
570 *Cognition*, 24:103–109.
- 571 [17] Kahana, M. J. (2012). *Foundations of human memory*. Oxford University Press, New
572 York, NY.
- 573 [18] Kimball, D. R., Smith, T. A., and Kahana, M. J. (2007). The fSAM model of false recall.
574 *Psychological Review*, 114(4):954–993.
- 575 [19] Manning, J. R. (2020). Context reinstatement. In Kahana, M. J. and Wagner, A. D.,
576 editors, *Handbook of Human Memory*. Oxford University Press.
- 577 [20] Manning, J. R., Hulbert, J. C., Williams, J., Piloto, L., Sahakyan, L., and Norman, K. A.
578 (2016). A neural signature of contextually mediated intentional forgetting. *Psychonomic*
579 *Bulletin and Review*, 23(5):1534–1542.

- 580 [21] Manning, J. R. and Kahana, M. J. (2012). Interpreting semantic clustering effects in
581 free recall. *Memory*, 20(5):511–517.
- 582 [22] Manning, J. R., Norman, K. A., and Kahana, M. J. (2015). The role of context in
583 episodic memory. In Gazzaniga, M., editor, *The Cognitive Neurosciences*, pages 557–566.
584 MIT Press.
- 585 [23] Manning, J. R., Polyn, S. M., Baltuch, G., Litt, B., and Kahana, M. J. (2011). Oscil-
586 latory patterns in temporal lobe reveal context reinstatement during memory search.
587 *Proceedings of the National Academy of Sciences, USA*, 108(31):12893–12897.
- 588 [24] Manning, J. R., Sperling, M. R., Sharan, A., Rosenberg, E. A., and Kahana, M. J.
589 (2012). Spontaneously reactivated patterns in frontal and temporal lobe predict semantic
590 clustering during memory search. *The Journal of Neuroscience*, 32(26):8871–8878.
- 591 [25] Masís-Obando, R., Norman, K. A., and Baldassano, C. (2022). Scheme representations
592 in distinct brain networks support narrative memory during encoding and retrieval.
593 *eLife*, 11:e70445.
- 594 [26] Momennejad, I., Russek, E. M., Cheong, J. H., Botvinick, M. M., Daw, N. D., and
595 Gershman, S. J. (2017). The successor representation in human reinforcement learning.
596 *Nature Human Behavior*, 1:680–692.
- 597 [27] Murdock, B. B. (1962). The serial position effect of free recall. *Journal of Experimental*
598 *Psychology: General*, 64:482–488.
- 599 [28] Nastase, S. A., Goldstein, A., and Hasson, U. (2020). Keep it real: rethinking the
600 primacy of experimental control in cognitive neuroscience. *NeuroImage*, 15(222):117254–
601 117261.

- 602 [29] Polyn, S. M. and Kahana, M. J. (2008). Memory search and the neural representation
603 of context. *Trends in Cognitive Sciences*, 12:24–30.
- 604 [30] Polyn, S. M., Norman, K. A., and Kahana, M. J. (2009). Task context and organization
605 in free recall. *Neuropsychologia*, 47:2158–2163.
- 606 [31] Postman, L. and Phillips, L. W. (1965). Short-term temporal changes in free recall.
607 *Quarterly Journal of Experimental Psychology*, 17:132–138.
- 608 [32] Raaijmakers, J. G. W. and Shiffrin, R. M. (1980). SAM: A theory of probabilistic search
609 of associative memory. In Bower, G. H., editor, *The Psychology of Learning and Motivation: Advances in Research and Theory*, volume 14, pages 207–262. Academic Press, New York,
610 NY.
611
- 612 [33] Radvansky, G. A. and Copeland, D. E. (2006). Walking through doorways causes
613 forgetting: situation models and experienced space. *Memory and Cognition*, 34(5):1150–
614 1156.
- 615 [34] Ranganath, C. and Ritchey, M. (2012). Two cortical systems for memory-guided
616 behavior. *Nature Reviews Neuroscience*, 13:713–726.
- 617 [35] Romney, A. K., Brewer, D. D., and Batchelder, W. H. (1993). Predicting clustering
618 from semantic structure. *Psychological Science*, 4:28–34.
- 619 [36] Schapiro, A. and Turk-Browne, N. (2015). Statistical learning. *Brain Mapping: An*
620 *Encyclopedic Reference*, 3:501–506.
- 621 [37] Sederberg, P. B., Howard, M. W., and Kahana, M. J. (2008). A context-based theory of
622 recency and contiguity in free recall. *Psychological Review*, 115(4):893–912.

- 623 [38] Shankar, K. H. and Howard, M. W. (2012). A scale-invariant internal representation
624 of time. *Neural Computation*, 24:134–193.
- 625 [39] Sirotin, Y. B., Kimball, D. R., and Kahana, M. J. (2005). Going beyond a single list:
626 modeling the effects of prior experience on episodic free recall. *Psychonomic Bulletin and*
627 *Review*, 12(5):787–805.
- 628 [40] Smith, S. M. and Vela, E. (2001). Environmental context-dependent memory: a review
629 and meta-analysis. *Psychonomic Bulletin and Review*, 8(2):203–220.
- 630 [41] Swallow, K. M., Barch, D. M., Head, D., Maley, C. J., Holder, D., and Zacks, J. M.
631 (2011). Changes in events alter how people remember recent information. *Journal of*
632 *Cognitive Neuroscience*, 23(5):1052–1064.
- 633 [42] Swallow, K. M., Zacks, J. M., and Abrams, R. A. (2009). Event boundaries in perception
634 affect memory encoding and updating. *Journal of Experimental Psychology: General*,
635 138(2):236–257.
- 636 [43] Welch, G. B. and Burnett, C. T. (1924). Is primacy a factor in association-formation.
637 *American Journal of Psychology*, 35:396–401.
- 638 [44] Xu, X., Zhu, Z., and Manning, J. R. (2022). The psychological arrow of time drives
639 temporal asymmetries in retrodicting versus predicting narrative events. *PsyArXiv*,
640 page doi.org/10.31234/osf.io/yp2qu.
- 641 [45] Yeshurun, Y., Swanson, S., Simony, E., Chen, J., Lazaridi, C., Honey, C. J., and
642 Hasson, U. (2017). Same story, different story: the neural representation of interpretive
643 frameworks. *Psychological Science*, 28(3):307–319.
- 644 [46] Ziman, K., Heusser, A. C., Fitzpatrick, P. C., Field, C. E., and Manning, J. R. (2018).

- 645 Is automatic speech-to-text transcription ready for use in psychological experiments?
646 *Behavior Research Methods*, 50:2597–2605.
- 647 [47] Zwaan, R. A., Langston, M. C., and Graesser, A. C. (1995). The construction of
648 situation models in narrative comprehension: an event-indexing model. *Psychological*
649 *Science*, 6(5):292–297.
- 650 [48] Zwaan, R. A. and Radvansky, G. A. (1998). Situation models in language compre-
651 hension and memory. *Psychological Bulletin*, 123(2):162–185.